

Visibility

- in the real world, opaque objects block light.
- we need to model this computationally
- one idea is to render back to front and use overwriting
 - this will have problem with visibility cycles
- we could explicitly store everything hit along a ray and then compute the closest
 - makes sense in a ray tracing setting, where we are working one pixel/ray at a time, but not for OpenGL, where we are working one triangle at a time.

z-buffer

- we will use use z-buffer
- triangles are drawn in any order
- each pixel in framebuffer stores “depth” value of closest geometry observed so far
- When a new triangle tries to set the color of a pixel, we first compare its depth to the value stored in the z-buffer. Only if the observed point in this triangle is closer do we overwrite the color and depth values of this pixel.
- this is done per-pixel, so no cycle problems.
- there are a optimizations where z-testing is done before the fragment shading is done

Other Uses of Visibility Calculations

- visibility to a light source is useful for shadows
 - we will talk about shadow mapping later
 - we will do shadow calculations in a ray tracer
- Visibility computation can also be used to speed up the rendering process.
 - If we know that some object is occluded from the camera, then we don't have to render the object in the first place.
 - can use a conservative test

Basic Mathematical Model

- for every point we define its $[x_n, y_n, z_n]^t$ coordinates using the following matrix expression.

$$\begin{bmatrix} x_n w_n \\ y_n w_n \\ z_n w_n \\ w_n \end{bmatrix} = \begin{bmatrix} x_c \\ y_c \\ z_c \\ w_c \end{bmatrix} = \begin{bmatrix} s_x & 0 & -c_x & 0 \\ 0 & s_y & -c_y & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix} \quad (1)$$

- we now also have the value $z_n = \frac{-1}{z_e}$.
- Our plan is to use this z_n value to do depth comparisons in our z-buffer.

correct ordering

- Given two points \tilde{p}^1 and \tilde{p}^2 with eye coordinates $[x_e^1, y_e^1, z_e^1, 1]^t$ and $[x_e^2, y_e^2, z_e^2, 1]^t$.
- Suppose that they both are in front of the eye, i.e., $z_e^1 < 0$ and $z_e^2 < 0$.
- And suppose that \tilde{p}^1 is closer to the eye than \tilde{p}^2 , that is $z_e^2 < z_e^1$.
- Then $-\frac{1}{z_e^2} < -\frac{1}{z_e^1}$, meaning $z_n^2 < z_n^1$.

projective transform

- we can now think of the process of taking points given by eye coordinates to points given by normalized device coordinates as an honest to goodness 3D geometric transformation.
- This kind of transformation is generally neither linear nor affine, but is something called a 3D *projective transformation*.
- projective transformations preserve co-linearity and co-planarity of points

projective figure

- first map film plane
 - iso- z_e so iso z_n
- map the red segment
 - some straight segment
- then note that rays must hit same pixel, so map to parallel lines

z_n interp is right

- preservation of coplanarity: for points on a fixed triangle, we will have $z_n = ax_n + by_n + c$, for some fixed a , b and c .
- Thus, the correct z_n value for a point can be computed using linear interpolation over the 2D image domain as long as we know its value at the three vertices of the triangle
- (see fig): projective transforms are funny
- linear interpolation of z_e values over the screen would produce wrong answer.
- “red” should win for entire bottom half of image.

Numerics

- there can be numerical difficulties when computing z_n . As z_e goes towards zero, the z_n value diverges off towards infinity.
- Conversely, points very far from the eye have z_n values very close to zero. The z_n of two such far away points may be indistinguishable in a finite precision representation, and thus the z-buffer will be ineffective in distinguishing which is closer to the eye.

solution: near/far

- solution: replacing the third row of the matrix with the more general row $[0, 0, \alpha, \beta]$.
- it is easy to verify that if the values α and β are both positive, then the z-ordering of points (assuming they all have negative z_e values) is preserved under the projective transform.
- To set α and β , we first select depth values n and f called the *near* and *far* values (both negative), such that our main region of interest in the scene is sandwiched between $z_e = n$ and $z_e = f$.
- Given these selections, we set $\alpha = \frac{f+n}{f-n}$ and $\beta = -\frac{2fn}{f-n}$.
- We can verify now that any point with $z_e = f$ maps to a point with $z_n = -1$ and that a point with $z_e = n$ maps to a point with $z_n = 1$
- Any geometry not in this [near..far] range is clipped away by OpenGL and ignored
- see fig

Code

- In OpenGL, use of the z-buffer is turned on with a call to `glEnable(GL_DEPTH_TEST)`.
- We also need a call to `glDepthFunc(GL_GREATER)`, since we are using a right handed coordinate system where “more-negative” is “farther from the eye”.
- in real life, you may see other conventions (for how to interpret n and f , some of the signs of the matrix, and the handedness of the ultimate z-test).