

path from vertex to pixel

- see pipeline fig
- three vertices have passed through the vertex shader
- follow its journey to be a bunch of pixels
- we might imagine that we need a divide-by-w step to position the vertices on the screen.
- we also need a few more steps.

Clipping

- processing for triangles that are fully or partially out of view.
- we don't want to see behind us
- we want to minimize processing
- the tricky part will be to deal with eye-spanning triangles.

eye spanners

- see figure
- back vertex projects higher up in the image
- filling in the in-between pixels will fill in the wrong region.
- solution: slice up the geometry by the six faces of the view frustum

coordinates for clipping

- if you wait for NDCs the vertex has flipped, and it's too late to do the clipping.
- could do in eye space, but then would need to use the camera parameters
- canonical solution: use clip coordinates, post matrix multiply but pre divide.
 - no divide = no flipping
- recall that we want points in the range

$$-1 < x_n < 1 \tag{1}$$

$$-1 < y_n < 1 \tag{2}$$

$$-1 < z_n < 1 \tag{3}$$

- in clip coordinates this is:

$$-w_c < x_c < w_c$$

$$-w_c < y_c < w_c$$

$$-w_c < z_c < w_c$$

divide

- clipping is done, we can now divide by w to obtain normalized device coordinates!

Backface Culling

- when drawing a closed solid object, we will only ever see one “front” side of each triangle.
- for efficiency we can drop these from the processing
- To do this, in OpenGL, we use the convention of ordering the three vertices in the draw call (IBO/VBO) so that they are counter clockwise (CCW) when looking at its front side.

- during setup, we call `glEnable(GL_CULL_FACE)`.
- to implement culling, OpenGL does the following:
- Let \tilde{p}_1 , \tilde{p}_2 , and \tilde{p}_3 be the three vertices of the triangle projected down to the $(x_n, y_n, 0)$ plane.
- Define the vectors $\vec{a} = \tilde{p}_3 - \tilde{p}_2$ and $\vec{b} = \tilde{p}_1 - \tilde{p}_2$.
- Next compute the cross product $\vec{c} = \vec{a} \times \vec{b}$.
- If the three vertices are counter clockwise in the plane, then \vec{c} will be in the z_n direction. Otherwise it will be in the positive $-z_n$ direction.
- When all the dust settles, this coordinate is

$$(x_n^3 - x_n^2)(y_n^1 - y_n^2) - (y_n^3 - y_n^2)(x_n^1 - x_n^2) \quad (4)$$

Viewport

- now we want to position the vertices in the window. so it is time to move to window coordinates
- each pixel center has an integer coordinate.
 - this will make subsequent pixel computations more natural.
- We want the lower left pixel center to have 2D window coordinates of $[0, 0]^t$ and the upper right pixel center to have coordinates $[W - 1, H - 1]^t$.
- We think of each pixel as owning the real estate which extends .5 pixel units in the positive and negative, horizontal and vertical directions from the pixel center.
- Thus the extent of 2D window rectangle covered by the union of all our pixels is the rectangle in window coordinates with lower left corner $[-.5, -.5]^t$ and upper right corner $[W - .5, H - .5]^t$.

viewport matrix

- we need a transform that maps the lower left corner to $[-.5, -.5]^t$ and upper right corner to $[W - .5, H - .5]^t$.
- the appropriate scale and shift can be done using the viewport matrix

$$\begin{bmatrix} x_w \\ y_w \\ z_w \\ 1 \end{bmatrix} = \begin{bmatrix} W/2 & 0 & 0 & (W-1)/2 \\ 0 & H/2 & 0 & (H-1)/2 \\ 0 & 0 & 1/2 & 1/2 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_n \\ y_n \\ z_n \\ 1 \end{bmatrix} \quad (5)$$

- this does a scale and shift in both x and y.
- you can verify that it maps the corners appropriately.
- In OpenGL, we set up this viewport matrix with the call `glViewport(0,0,W,H)`.
- The third row of this matrix is used to map the $[-1..1]$ range of z_n values to the more convenient $[0..1]$ range.
- so now (in our conventions), $z_w = 0$ is far and $z_w = 1$ is near.
 - so we must also tell OpenGL that when we clear the z-buffer, we should set it to 0; we do this with the call `glClearDepth(0.0)`

texture Viewport

- the abstract domain for textures is not the canonical square, but instead is the *unit square*
- in this case the coordinate transformation matrix is

$$\begin{bmatrix} x_w \\ y_w \\ - \\ 1 \end{bmatrix} = \begin{bmatrix} W & 0 & 0 & -1/2 \\ 0 & H & 0 & -1/2 \\ - & - & - & - \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_t \\ y_t \\ - \\ 1 \end{bmatrix} \quad (6)$$

Rasterization

- Starting from the window coordinates for the three vertices, the rasterizer needs to figure out which pixel-centers are inside of the triangle.
- Each triangle on the screen can be defined as the intersection of three half-spaces.
- Each such halfspace is defined by a line that coincides with one of the edges of the triangle, and can be tested using an “edge function” of the form

$$\text{edge} = ax_w + by_w + c$$

where the (a, b, c) are constants that depend on the geometry of the edge.

- A positive value of this function at a pixel with coordinates $[x_w, y_w]^t$ means that the pixel is inside the specified halfspace.
- If all three tests pass, then the pixel is inside the triangle.

speed up

- only look at pixels in the bounding box of the triangle
- test if a pixel block is entirely outside of triangle
- use incremental calculations along a scanline

interpolation

- As input to rasterization, each vertex also has some auxiliary data associated with it.
 - This data includes a z_w value,
 - as well as other data that is related, but not identical to the varying variables
- It is also the job of the rasterizer to linearly interpolate this data over the triangle.
- Each such value v to be linearly interpolated can be represented as an affine function over screen space with the form

$$v = ax_w + by_w + c \tag{7}$$

- An affine function can be easily evaluated at each pixel by the rasterizer.
- Indeed, this is no different from evaluating the edge test functions just described.

boundaries

- for pixel on edge or vertex it should be rendered exactly once
- need special care in the implementation.