

Computer Science 175

Introduction to Computer Graphics

www.fas.harvard.edu/~lib175

time: m/w 2:30-4:00 pm

place:p301

section times: tba

Instructor: Steven “shlomo” Gortler

www.cs.harvard.edu/~sjg

sjg@cs.harvard.edu

office: md 243

teaching assistants:

Yuanchen Zhu

yzhu@fas.harvard.edu

Fangyang Shen

cs175@shenfy.com

textbook

- “Foundations of 3D Computer Graphics”, S.J. Gortler
 - Not required, but is a good mirror of the course.
- web has resources on OpenGL/GLSL syntax

prerequisites

- ability to program in C (multi-hundred line program)
 - we will use C++, but will help get you up to speed
- some familiarity with linear algebra

warnings

- you will need to use computers that support programmable shaders
- most of the assignments will build off each other
- much of the openGL gobbledygook will be given as a black box.
- I expect to interact in class with the students, so I expect students to be engaged. (no laptops, smartphones,...) (fig)
- I don’t capitalize or spell well.

graded work

- 10 programming assignments
 - a few written parts
 - one weekers
 - due (mostly) at mondays at 11:59pm
- unit per program
- 1 final project
 - 1.5 units
- participation
 - 1 unit or so

late policy

- programs due on midnight
- you have 5 free days to use as you wish
- up to one day late: -15%
- up to two days late: -30%
- after two days late: 0%
- we will optimize late penalties/free days at the end.

Standards

- you can work/submit in pairs if you wish
 - you can change this arrangement for each asst.
- you should not look at code from other groups
- you should not discuss written assts with other groups
- we will use Piazza

platforms for class

- science center linux boxes
- windows
 - visual studio 2010
 - * to get a free copy send email to holloway@eecs.harvard.edu
 - cygwin/gcc if you can get it to work
- mac
- must have recent generation graphics cards that support vertex and fragment shaders
 - first assignment will get this worked out

lets look at the assignments

- 1: hello world openGL
- 2-3: 3d viewer
- 4: robots and picking
- 5-6: keyframe animator
- 7: meshes and subdivision
- 8: fur
- 9: under the hood: texture coordinate interpolation and reconstruction
- (10: ray tracer)
- 11: final project

applications:

- entertainment
 - video games
 - special effects

- animation
- virtual world interactions
- computer aided design
 - car bodies
 - architectural environments
- simulation
 - flight simulators
 - surgical simulators
- 3D visualization
 - medical scan imagery
 - fluid flow

syllabus

- we understand how to work on top of OpenGL
- we will understand what is happening inside of OpenGL

topics

- getting started with OpenGL
- coordinate systems and transforms
- quaternions
- interpolation and splines for key frame animation
- color theory
- overview of more advanced geometric modeling
- overview of more advanced computer animation
- subdivision surfaces
- camera projections
- rasterization process
- image reconstruction and filtering (dots)
- shading and shadows
- ray tracing

class will follow the real time pipeline

- basic representation of object starts with a collection of triangles (fig) (demo)
- each triangle is described by 3 vertices
- each vertex is described by (x,y,z) coordinates
- other information is attached to each vertex
 - color, normal vector
- this data is called *attribute* data.
- this data is passed (once) to OpenGL and stored in a *vertex buffer*

- later we can make an OpenGL *draw* call

vertex processing

- after a draw call
- each vertex (set of attributes) is passed through a *vertex shader* that you write and load into OpenGL.
- shader also has access to *uniform variables* that you set
- typically does “geometric” transformations on the vertex coordinates
- to place the objects correctly with respect to each other and the eye
- to get perspective effect
- output position as `gl_Position` as well as user defined *varying variables*

fixed function

- 3 vertices for each triangle are collected by the *assembler*
- triangles vertices are positioned within the window
- triangles are *rasterized*
 - which pixels are inside the triangle
 - varying variable data (color, normal, texture coordinates) are appropriately *interpolated* for each pixel

pixel processing

- information for each dot is sent through a *fragment shader* that you write and load into OpenGL
 - figures out final output color of the pixel from the interpolated data
- outputs a color for the screen
- typically simulates material reflection (demo)
- also can look up data in a *texture* image (fig)
 - pointers into the texture are done using varying variables called *texture coordinates*
 - adds visual complexity

merging

- decide whether to write into framebuffer
- looks at data already there
- main thing: depth test. (demo)
- also alpha blending (transparency)