

camera transforms

- until now we have considered all of our geometry in a 3d space
- ultimately everything ended up in eye coordinates with coordinates $[x_e, y_e, z_e, 1]^t$.
- we said that the camera is placed at the origin of the eye frame \vec{e}^t , and that it is looking down the eye's negative z-axis.
- this somehow produces a 2d image.
- we had a magic matrix which created `gl_Position`
- now we will study this step.

Pinhole Camera model

- see fig
- As light travels towards the film plane, most is blocked by an opaque surface placed at the $z_e = 0$ plane.
- But we place a very small hole in the center of the surface, at the point with eye coordinates $[0, 0, 0, 1]^t$.
- Only rays of light that pass through this point reach the film plane and have their intensity recorded on film. The image is recorded at a film plane placed at, say, $z_e = 1$.
 - a physical camera needs a finite aperture and a lens, but we will ignore this.
- to avoid the image flip, we can mathematically model this with the film plane *in front* of the pinhole, say at the $z_e = -1$
- if we hold up the photograph at the $z_e = -1$ plane, and observe it with our own eye placed at the origin (see Figure), it will look to us just like the original scene would have.

Basic Mathematical Model

- let \tilde{p} have eye coordinates $[x_e, y_e, z_e]^t$
- where does the ray from \tilde{p} to the origin hits the film plane?
- all points on the ray hit the same pixel.
- all points on the ray are all scales of each other
- Let us use eye coordinates $[x_n, y_n, -1]^t$ to specify the hit point on our film plane.
- all points on ray must be of the form : $\alpha[x_n, y_n, -1]^t$
- so $[x_e, y_e, z_e]^t$ is of this form with $\alpha = -z_e$.
- so $[x_e, y_e]^t = -z_e[x_n, y_n]^t$
- so

$$x_n = -\frac{x_e}{z_e} \quad (1)$$

$$y_n = -\frac{y_e}{z_e} \quad (2)$$

$$(3)$$

in matrix form

- We can model this expression as a matrix operation as follows.

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix} = \begin{bmatrix} x_c \\ y_c \\ - \\ w_c \end{bmatrix} = \begin{bmatrix} x_n w_n \\ y_n w_n \\ - \\ w_n \end{bmatrix} \quad (4)$$

- The raw output of the matrix multiply, $[x_c, y_c, -, w_c]^t$, are called the *clip coordinates* of \tilde{p} .

- $w_n = w_c$ is a new variable called the w-coordinate.
- In such clip coordinates, the fourth entry of the coordinate 4-vector is not necessarily a zero or a one.

divide by w

- We say that $x_n w_n = x_c$ and $y_n w_n = y_c$. If we want to extract x_n alone, we must perform the division $x_n = \frac{x_n w_n}{w_n}$
- this recovers our camera model
- Our output coordinates, with subscripts “n”, are called *normalized device coordinates* because they address points on the image in abstract units without specific reference to numbers of pixels.
- we keep all of the image data in the *canonical square* $-1 \leq x_n \leq +1$, $-1 \leq y_n \leq +1$, and ultimately map this onto a window on the screen.
 - Data outside of this square is not be recorded or displayed.
 - This is exactly the model we used to describe 2D OpenGL

scales

- By changing the entries in the projection matrix we can slightly alter the geometry of the camera transformation.
- we could push the film plane out to $z_e = n$, where n is some negative number (zoom lens)
- eye coordinates of points on the film are of the form: $[x_n, y_n, n]^t$
- points along its ray are of the form $\alpha[x_n, y_n, n]^t$
- so a point $[x_e, y_e, z_e]^t$ on the ray is of this form with $\alpha = \frac{z_e}{n}$.
- so a point $[x_e, y_e, z_e]^t$ on the ray must satisfy $[x_e, y_e]^t = \frac{z_e}{n}[x_n, y_n]^t$
- and

$$\begin{aligned} x_n &= \frac{x_e n}{z_e} \\ y_n &= \frac{y_e n}{z_e} \end{aligned}$$

in matrix form

- In matrix form, this becomes

$$\begin{bmatrix} x_n w_n \\ y_n w_n \\ - \\ w_n \end{bmatrix} = \begin{bmatrix} -n & 0 & 0 & 0 \\ 0 & -n & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix}$$

- note this matrix is the same as

$$\begin{bmatrix} -n & 0 & 0 & 0 \\ 0 & -n & 0 & 0 \\ - & - & - & - \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

- this has the same effect as starting with our original camera, scaling by $-n$, and cropping to the canonical square.
- so at this point, it we may as well think of the $[x_n, y_n]^t$ as coordinates resulting from a camera transform (and not the eye coordinates of points on a film plane).
 - so now we can simply think up some constraints on this transform and find the correct matrix.

fovY

- scale can be determined by vertical angular *field of view* of the desired camera.

- if we want our camera to have a field of view of θ degrees.
- we want any point with the maximal vertical angle to map to the boundary of the image.
- one such point has eye coordinates: $[0, \tan(\frac{\theta}{2}), -1, 1]^t$
- we want the point with eye coordinates: $[0, \tan(\frac{\theta}{2}), -1, 1]^t$ maps to normalized device coordinates $[0, 1]^t$.
- so we can use the matrix

$$\begin{bmatrix} \frac{1}{\tan(\frac{\theta}{2})} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan(\frac{\theta}{2})} & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \quad (5)$$

dealing with aspect ratio

- Suppose the window is wider than it is high. In our camera transform, we need to squish things horizontally so a wider horizontal field of view fits into our retained canonical square.
- When the data is later mapped to the window, it will be stretched out correspondingly and will not appear distorted.
- Define a , the *aspect ratio* of a window, to be its width divided by its height (measured say in pixels).
- We can then set our projection matrix to be

$$\begin{bmatrix} \frac{1}{a \tan(\frac{\theta}{2})} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan(\frac{\theta}{2})} & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \quad (6)$$

- so when the window is wide, we will keep more horizontal FOV, and when the window is tall, we will keep less horizontal FOV
- as an alternative, we could have an `fovMin`, and when the window is tall, we would need to calculate an appropriate larger `fovY`.
 - and then build the matrix.
- this is what we do in our code.

FOV issues

- to be a “window” onto the world, the FOV should match the angular extents of the window in the viewers field.
- this might give a too limited view onto the world
- so we can increase it to see more
- but this might give a somewhat unnatural look (demo)

motivate shifts

- look at the two street pics
 - what is the difference between the two cameras.
 - hint: if a geometric plane recedes from the film, it appears smaller.
- imagine we want to model the screen as a window, what should this camera look like

Shifts

- sometimes, we wish to crop the image non-centrally

- this can be modeled as translating the NDC's and then cropping centrally.

$$\begin{aligned} \begin{bmatrix} x_n w_n \\ y_n w_n \\ - \\ w_n \end{bmatrix} &= \begin{bmatrix} 1 & 0 & 0 & c_x \\ 0 & 1 & 0 & c_y \\ - & - & - & - \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix} \\ &= \begin{bmatrix} 1 & 0 & -c_x & 0 \\ 0 & 1 & -c_y & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix} \end{aligned}$$

- also useful for tiled displays, stereo viewing.

frustum

- shifts are often specified by first specifying a near plane $z_e = n$.
- On this plane, a rectangle is specified with the eye coordinates of an axis aligned rectangle. (For non-distorted output, the aspect ratio of this rectangle should match that of the final window.)
 - using l, r, t, b .

$$\begin{bmatrix} -\frac{2n}{r-l} & 0 & \frac{r+l}{r-l} & 0 \\ 0 & -\frac{2n}{t-b} & \frac{t+b}{t-b} & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \quad (7)$$

- (fig)

Context

- projection could be applied to every point in the scene
- in CG, we will apply it to the verts to position a triangle on the screen
- the rest of the triangle will then get filled in on the screen as we shall see.